

[Sign up](#)

pingcap / awesome-database-learning Public

A list of learning materials to understand databases internals

☆ 5.5k stars 🍴 723 forks

☆ Star

🔔 Notifications

[Code](#) [Issues](#) 2 [Pull requests](#) 3 [Security](#) [Insights](#)

🔗 master ▾



SahilTalrejacoder Added a new resource (#51) ...

Sep 2, 2022 ⌚ 48

[View code](#)

☰ README.md

🔗 Awesome Database Learning

A list of learning materials to understand databases internals, including but not limited to:

- papers
- blogs
- courses
- talks

Please submit a pull request if there is any material that you think should be included in this collection.

🔗 Table of Contents

- [Recommended Courses, Books and Talks](#)
 - [Courses](#)
 - [Books](#)
 - [Talks](#)
 - [Blogs](#)
- [SQL & Relation Algebra](#)

- Query Optimizer
 - Planner Models
 - Subquery Optimization
 - Join Order Optimization
 - Functional Dependency & Physical Properties
 - Cost Model
 - Statistics
- Query Execution
 - Execution Framework
 - Vectorization vs Compilization
 - Join
 - Hash Table
 - Bloom Filter
- DDL
- Relational Model
 - Codd's Rules
 - Relational Data Model
 - Relational Algebra
 - ER to Relational Model
 - SQL - Overview
- Transaction
 - Isolation Levels
 - Concurrency Control
- Network
- Storage
 - NoSQL Systems
 - Buffer Management
 - Disk IO
 - B-Tree
 - LSM-Tree
 - Learned Indexes Structures
- Serializing & RPC
- Data Partitioning
- Replication & Consistency
- Consensus
- Scheduling
- Benchmark & Testing
- HTAP
- TLA+

🔗 Recommended Courses, Books and Talks

🔗 Courses

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
- CMU [Advanced Database Systems \(15-721\)](#), thanks to [Andy Pavlo](#)
- UC Berkeley [Introduction to Database Systems](#)
- Stanford [Database System Implementation](#)
- Cornell [Introduction to Database Systems](#) by Prof. Trummer
- [Let's Build a Simple Database](#), thanks to [cstack](#)

🔗 Books

- Stanford [Database Systems: The Complete Book](#)
- [Designing Data-Intensive Applications](#), [中文翻译](#)
- [Database Internals](#)
- [Foundations of Databases](#)
- [Readings in Database Systems](#), 5th Edition
- [Database Design and Implementation: Second Edition \(Data-Centric Systems and Applications\)](#)
- [Principles of Distributed Database Systems](#), 4th ed
- [Inside SQLite](#)
- [Architecture of a Database System](#)
- [Relational Database Index Design and the Optimizers](#)
- [Transactional Information Systems: Theory, Algorithms, and the Practice of Concurrency Control](#)

🔗 Talks

- [Data Structures and Algorithms for Big Databases](#)
- [A Journey From A Quick Hack To A High-Reliability Database Engine](#)

🔗 Blogs

- [How does a relational database work](#)
- [The Internals of PostgreSQL](#)
- [Books propose](#)
- [what is database and its types](#)

🔗 SQL & Relation Algebra

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Course Introduction and the Relational Model](#)
 - [Advanced SQL](#)
- UC Berkeley [Introduction to Database Systems](#)
 - Introduction + SQL I
 - SQL II
 - Relational Algebra

🔗 Query Optimizer

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Query Planning & Optimization I](#)
 - [Query Planning & Optimization II](#)

Blogs:

- [数据库内核杂谈](#), thanks to [顾仲贤](#)
 - [数据库内核杂谈（七）：数据库优化器（上）](#)
 - [数据库内核杂谈（八）：数据库优化器（下）](#)
- [SQL优化器原理 - 查询优化器综述](#), thanks to [勿烦](#)

🔗 Planner Models

Blogs:

- [数据库内核杂谈](#), thanks to [顾仲贤](#)
 - [数据库内核杂谈（九）：开源优化器 ORCA](#)
- [SQL 查询优化原理与 Volcano Optimizer 介绍](#), thanks to [张茄子](#)
- [Cascades Optimizer](#), thanks to [hellocode](#)

Papers:

- 1979, [Access Path Selection in a Relational Database Management System](#), SIGMOD
- 1979, [Query Processing in Main Memory Database Management Systems](#), VLDB
- 1987, [Query Optimization by Simulated Annealing](#), SIGMOD
- 1988, [Grammar-like Functional Rules for Representing Query Optimization Alternatives](#), SIGMOD
- 1993, [The Volcano Optimizer Generator- Extensibility and Efficient Search](#), ICDE
- 1995, [The Cascades Framework for Query Optimization](#), IEEE Data engineering

Bulltin

- 1998, [An Overview of Query Optimization in Relational Systems](#), PODS
- 2001, [LEO – DB2's LEarning Optimizer](#), VLDB
- 2004, [Robust Query Processing through Progressive Optimization](#), SIGMOD
- 2014, [Orca: A Modular Query Optimizer Architecture for Big Data](#), SIGMOD
- 2016, [Parallelizing Query Optimization on Shared-Nothing Architectures](#), VLDB
- 2016, [The MemSQL Query Optimizer: A modern optimizer for real-time analytics in a distributed database](#), VLDB

🔗 Subquery Optimization

Blogs:

- [SQL 子查询的优化](#), thanks to [Eric Fu](#)
- [Calcite 子查询处理 - I \(RemoveSubQuery\)](#), thanks to [一只无情的小猫咪](#)
- [Calcite 子查询处理 - II \(Decorrelate\)](#), thanks to [一只无情的小猫咪](#)

Papers:

- 2001, [Orthogonal Optimization of Subqueries and Aggregation](#), SIGMOD
- 2009, [Enhanced subquery optimizations in Oracle](#), VLDB
- 2015, [Unnesting Arbitrary Queries](#), BTW

🔗 Join Order Optimization

Papers:

- 2006, [Analysis of Two Existing and One New Dynamic Programming Algorithm for the Generation of Optimal Bushy Join Trees without Cross Products](#), VLDB
- 2015, [How Good Are Query Optimizers, Really?](#), VLDB
- 2018, [Adaptive Optimization of Very Large Join Queries](#), SIGMOD

🔗 Functional Dependency & Physical Properties

Thesis:

- 2000, [Exploiting Functional Dependence in Query Optimization](#)

Papers:

- 1996, [Fundamental Techniques for Order Optimization](#), SIGMOD
- 2004, [An Efficient Framework for Order Optimization](#), ICDE
- 2010, [Incorporating Partitioning and Parallel Plans into the SCOPE Optimizer](#), ICDE

🔗 Cost Model

Papers:

- 1996, [Modelling Costs for a MM-DBMS](#), in Real-Time Databases
- 2014, [Approximation Schemes for Many-Objective Query Optimization](#), SIGMOD
- 2015, [Multi-Objective Parametric Query Optimization](#), VLDB

🔗 Statistics

Papers:

- 1984, [Accurate Estimation of the Number of Tuples Satisfying a Condition](#), SIGMOD
- 1993, [Optimal Histograms for Limiting Worst-Case Error Propagation in the Size of Join Results](#), ACM Trans. on Database Systems
- 1993, [Universality of Serial Histograms](#), VLDB
- 1995, [Balancing Histogram Optimality and Practicality for Query Result Size Estimation](#), SIGMOD
- 1996, [Improved Histograms for Selectivity Estimation of Range Predicates](#), SIGMOD
- 1997, [SEEKING the truth about ad hoc join costs](#), VLDB
- 2000, [Towards Estimation Error Guarantees for Distinct Values](#), SIGMOD/PODS
- 2001, [Distinct Sampling for Highly-Accurate Answers to Distinct Values Queries and Event Reports](#), VLDB
- 2003, [The History of Histograms](#), VLDB
- 2005, [An Improved Data Stream Summary: The Count-Min Sketch and its Applications](#), Journal of Algorithms
- 2007, [New Estimation Algorithms for Streaming Data: Count-min Can Do More](#)
- 2009, [Preventing Bad Plans by Bounding the Impact of Cardinality Estimation Errors](#), VLDB
- 2010, [Histograms Reloaded: The Merits of Bucket Diversity](#), SIGMOD
- 2014, [Exploiting Ordered Dictionaries to Efficiently Construct Histograms with Q-Error Guarantees in SAP HANA](#), SIGMOD
- 2017, [Adaptive Statistics in Oracle 12c](#), VLDB
- 2019, [Pessimistic Cardinality Estimation: Tighter Upper Bounds for Intermediate Join Cardinalities](#), SIGMOD
- 2019, [Deep Unsupervised Cardinality Estimation](#), VLDB
- 2020, [NeuroCard: One Cardinality Estimator for All Tables](#), VLDB

Books:

- [Synopses for Massive Data: Samples, Histograms, Wavelets, Sketches](#)

🔗 Query Execution

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Query Execution I](#)
 - [Query Execution II](#)

🔗 Execution Framework

Papers:

- 1994, [Volcano-An Extensible and Parallel Query Evaluation System](#), IEEE Transactions on Knowledge and Data Engineering February
- 2014, [Morsel-Driven Parallelism: A NUMA-Aware Query Evaluation Framework for the Many-Core Age](#), SIGMOD

🔗 Vectorization vs Compilization

Blogs:

- [Overhead of a Generalized Query Execution Engine](#), from [The Pivotal Engineering Journal](#), thanks to the Pivotal Engineering team

Papers:

- 2005, [MonetDB/X100: Hyper-Pipelining Query Execution](#), CIDR
- 2011, [Efficiently Compiling Efficient Query Plans for Modern Hardware](#), VLDB
- 2017, [Relaxed Operator Fusion for In-Memory Databases: Making Compilation, Vectorization, and Prefetching Work Together At Last](#), VLDB
- 2018, [Everything You Always Wanted to Know About Compiled and Vectorized Queries But Were Afraid to Ask](#), VLDB
- 2018, [Adaptive Execution of Compiled Queries](#), ICDE

🔗 Join

Papers:

- 2013, [Multi-Core, Main-Memory Joins: Sort vs. Hash Revisited](#), VLDB
- 2017, [Looking Ahead Makes Query Plans Robust](#), VLDB

🔗 Hash Table

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Hash Tables](#)

Blogs:

- [Fibonacci Hashing: The Optimization that the World Forgot \(or: a Better Alternative to Integer Modulo\)](#), thanks to [Malte Skarupke](#)
- [All hash table sizes you will ever need](#), thanks to [Database Architects - Thomas Neumann](#)

🔗 Bloom Filter

Papers:

- 2018, [SuRF: Practical Range Query Filtering with Fast Succinct Tries](#), SIGMOD

🔗 DDL

- 2013, [Online, Asynchronous Schema Change in F1](#), VLDB

🔗 Relational Model

Blogs:

- [What is a Relational Database?](#), thanks to [Caleb Curry](#)
- [What is a Relational Database?](#), thank to [JAMES GALLAGHER](#)

🔗 Codd's Rules

Blogs:

- [Codd's Rules for Relational Database Systems](#), thanks to [Kevin Kline](#)

🔗 Relational Data Model

Blogs:

- [Relational model](#), thanks to [Wikipedia](#)

🔗 Relational Algebra

Blogs:

- [Introduction of Relational Algebra in DBMS](#), thanks to [GeeksforGeeks](#)

🔗 ER to Relational Model

Blogs:

- [ER Model to Relational Model](#), thanks to [tutorialspoint](#)

🔗 SQL - Overview

Blogs:

- [An Overview of SQL Text Functions](#), thanks to [Zahin Rahman](#)

🔗 Transaction

🔗 Isolation Levels

Blogs:

- [一致性模型](#), thanks to [siddontang](#)

Papers:

- 1995, [A Critique of ANSI SQL Isolation Levels](#), SIGMOD
- 2000, [Generalized Isolation Level Definitions](#), Proceedings of 16th International Conference on Data Engineering

🔗 Concurrency Control

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Concurrency Control Theory](#)
 - [Two-Phase Locking Concurrency Control](#)
 - [Timestamp Ordering Concurrency Control](#)
 - [Multi-Version Concurrency Control](#)
- CMU [Advanced Database Systems \(15-721\)](#), thanks to [Andy Pavlo](#)
 - [Multi-Version Concurrency Control \(Design Decisions\)](#)
 - [Multi-Version Concurrency Control \(Protocols\)](#)
 - [Multi-Version Concurrency Control \(Garbage Collection\)](#)

Papers:

- 1976, [The Notions of Consistency and Predicate Locks in a Database System](#), Communications of the ACM
- 1981, [Concurrency Control in Distributed Database Systems](#), ACM Computing Surveys
- 1981, [On Optimistic Methods for Concurrency Control](#), ACM Transactions on Database Systems
- 1983, [Multiversion Concurrency Control - Theory and Algorithms](#), ACM Transactions on Database Systems

- 2012, [Serializable Snapshot Isolation in PostgreSQL](#), VLDB
- 2012, [Calvin: Fast Distributed Transactions for Partitioned Database Systems](#), SIGMOD
- 2014, [MaaT: effective and scalable coordination of distributed transactions in the cloud](#), VLDB
- 2014, [Staring into the Abyss: An Evaluation of Concurrency Control with One Thousand Cores](#), VLDB
- 2014, [An Evaluation of the Advantages and Disadvantages of Deterministic Database Systems](#), VLDB
- 2015, [Fast Serializable Multi-Version Concurrency Control for Main-Memory Database Systems](#), SIGMOD
- 2017, [An Empirical Evaluation of In-Memory Multi-Version Concurrency Control](#), VLDB
- 2017, [An Evaluation of Distributed Concurrency Control](#), VLDB
- 2019, [Scalable Garbage Collection for In-Memory MVCC Systems](#), VLDB

🔗 Network

Courses:

- CMU [Advanced Database Systems \(15-721\)](#), thanks to [Andy Pavlo](#)
 - [Networking Protocols](#)

Papers:

- 2016, [The End of Slow Networks: It's Time for a Redesign](#), VLDB
- 2016, [Accelerating Relational Databases by Leveraging Remote Memory and RDMA](#), SIGMOD
- 2017, [Don't Hold My Data Hostage: A Case for Client Protocol Redesign](#), VLDB

🔗 Storage

🔗 NoSQL Systems

Papers:

- 2006, [Bigtable: A Distributed Storage System for Structured Data](#), OSDI
- 2007, [Dynamo: Amazon's Highly Available Key-value Store](#), SOSP
- 2008, [PNUTS: Yahoo!'s Hosted Data Serving Platform](#), VLDB
- 2010, [Cassandra - A Decentralized Structured Storage System](#), SOSP
- 2019, [PNUTS to Sherpa: Lessons from Yahoo!'s Cloud Database](#), VLDB

🔗 Buffer Management

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Buffer Pools](#)

Papers:

- 1987, [The 5 Minute Rule for Trading Memory for Disc Accesses and the 5 Byte Rule for Trading Memory for CPU Time](#), SIGMOD
- 2008, [The Five Minute Rule 20 Years Later and How Flash Memory Changes the Rules](#), ACM Queue
- 2018, [Managing Non-Volatile Memory in Database Systems](#), SIGMOD
- 2018, [LeanStore: In-Memory Data Management Beyond Main Memory](#), ICDE
- 2020, [Umbra: A Disk-Based System with In-Memory Performance](#), CIDR

🔗 Disk IO

Blogs:

- [On Disk IO, Part 1: Flavors of IO](#), thanks to [Alex](#)
- [On Disk IO, Part 2: More Flavours of IO](#), thanks to [Alex](#)
- [On Disk IO, Part 3: LSM Trees](#), thanks to [Alex](#)
- [On Disk IO, Part 4: B-Trees and RUM Conjecture](#), thanks to [Alex](#)
- [On Disk IO, Part 5: Access Patterns in LSM Trees](#), thanks to [Alex](#)
- [Ensuring data reaches disk\(LWN\)](#)
- [Read, write & space amplification - pick 2](#), thanks to [Mark Callaghan](#)

Papers:

- 2016, [Design Tradeoffs of Data Access Methods](#), SIGMOD
- 2016, [Designing Access Methods: The RUM Conjecture](#), EDBT

🔗 B-Tree

Blogs:

- [B树、B+树索引算法原理（上）](#) thanks to [codedump](#)
- [B树、B+树索引算法原理（下）](#)

Courses:

- CMU [Database Systems \(15-445/645\)](#), thanks to [Andy Pavlo](#)
 - [Trees Indexes I](#)
 - [Trees Indexes II](#)

- CMU [Advanced Database Systems \(15-721\)](#), thanks to [Andy Pavlo](#)
 - [OLTP Indexes \(B+Tree Data Structures\)](#)

Papers:

- 1979, [The Ubiquitous B-Tree](#)

↻ LSM-Tree

Papers:

- 1996, [The Log-Structured Merge-Tree \(LSM-Tree\)](#),
- 2014, [A Comparison of Fractal Trees to Log-Structured Merge \(LSM\) Trees](#)
- 2017, [WiscKey: Separating Keys from Values in SSD-conscious Storage](#), TOS
- 2019, [LSM-based Storage Techniques: A Survey](#)

↻ Learned Indexes Structures

Papers:

- 2018, [The Case for Learned Index Structures](#)
- 2019, [Learning Multi-dimensional Indexes](#)
- 2020, [XIndex: A Scalable Learned Index for Multicore Data Storage](#)
- 2020, [RadixSpline: A Single-Pass Learned Index](#), [Source Code](#), aiDM@SIGMOD
- 2020, [The PGM-index: a fully-dynamic compressed learned index with provable worst-case bounds](#), [Source Code](#), VLDB
- 2020, [From WiscKey to Bourbon: A Learned Index for Log-Structured Merge Trees](#)

↻ Serializing & RPC

- [Protocol Buffers Developer Guide](#)
- [gRPC Documentation](#)

↻ Data Partitioning

Blogs:

- [TiDB Internal \(I\) - Data Storage](#)
- [Partitioning Behavior of DynamoDB](#), thanks to [Parth Modi](#)

Papers:

- 2007, [Dynamo: Amazon's Highly Available Key-value Store](#), SOSP

↻ Replication & Consistency

Blogs:

- [Tick or Tock? Keeping Time and Order in Distributed Databases](#), thanks to [Liu Tang](#)

Papers:

- 2012, [Consistency Tradeoffs in Modern Distributed Database System Design](#)
- 2020, [Strong and Efficient Consistency with Consistency-Aware Durability](#), FAST 2020

↻ Consensus

Technical report:

- University of Cambridge [Distributed consensus revised](#), a great paper about Consensus especially Paxos and Paxos-Related algorithms, by Heidi Howard

Papers:

- 2014, [Ark: A Real-World Consensus Implementation](#), CoRR

↻ Scheduling

Blogs:

- [Building a Large-scale Distributed Storage System Based on Raft](#), by Ed Huang

Papers:

- 2016, [Automated Demand-driven Resource Scaling in Relational Database-as-a-Service](#), SIGMOD
- 2019, [Autoscaling Tiered Cloud Storage in Anna](#), VLDB
- 2020, [Adaptive HTAP through Elastic Resource Scheduling](#), SIGMOD
- 2020, [MorphoSys: Automatic Physical Design Metamorphosis for Distributed Database Systems](#), VLDB

↻ Benchmark & Testing

Blogs:

- [Use go-ycsb to benchmark different databases \(1\)](#), thanks to [siddontang](#)
- [Chaos Tools and Techniques for Testing the TiDB Distributed NewSQL Database](#), thanks to [Liu Tang](#)
- [Creating Custom Sysbench Scripts](#), thanks to [Matthew Boehm](#)

Papers:

- 2010, [Benchmarking Cloud Serving Systems with YCSB](#), SOCC

HTAP

Papers:

- 2020, [TiDB: A Raft-based HTAP Database](#), VLDB
- 2020, [F1 Lightning: HTAP as a Service](#), VLDB

TLA+

Talks:

- [The TLA+ Video Course](#)


Releases

No releases published

Packages

No packages published

Contributors 24



[+ 13 contributors](#)

[Terms](#) [Privacy](#) [Security](#) [Status](#) [Docs](#) [Contact GitHub](#) [Pricing](#) [API](#) [Training](#) [Blog](#) [About](#)



© 2022 GitHub, Inc.