



SNOWFLAKE FOR **DATA LAKE**



SNOWFLAKE'S PLATFORM




SNOWFLAKE FOR DATA LAKE

Exceptional Query Performance



Achieve the full potential of your data with accelerated analytics and unlimited scale

Integrated Data Pipelines with Extensibility



Run simple, but reliable and performant data pipelines that meet your unique needs

Governed & Secure Collaboration



Enable collaboration across your data lake



Complement Your Data
Lake with Snowflake

Build Your Data Lake
with Snowflake



Complement Your Data
Lake with Snowflake

Build Your Data Lake
with Snowflake



DATA LAKE WITH SNOWFLAKE

1

2

3



WORKFLOW



Data Ingestion




Data Transformation



Consumption



DATA LAKE WITH SNOWFLAKE



**SNOWFLAKE
CLOUD DATA
PLATFORM**



One platform,
one copy of data,
many workloads



Secure and
governed access
to all data



Near-zero
maintenance,
as a service



Unlimited
performance
and scale



DATA LAKE WITH SNOWFLAKE

CLICK ON NAV BAR TO GO TO DEEP DIVE

1

2

3



WORKFLOW



Data Ingestion



Data Transformation



Consumption

DATA LAKE
FEATURES



Directly query data in your data lake with **External Tables**



Better query performance using **Materialized Views** over External Tables



Synchronize External Tables with **Apache Hive** metastore



Automatically register new files with **auto-refresh of partitions**



Accelerate data exploration with **Snowsight**



Elastic Scaling with multi-clustered warehouse



Secure **data sharing & data exchange**



Encryption, RBAC, PCI, HIPAA, FedRAMP



Dynamic **Data Masking** & External **Tokenization**

PLATFORM
FEATURES

SNOWFLAKE
CLOUD DATA
PLATFORM



One platform,
one copy of data,
many workloads



Secure and
governed access
to all data



Near-zero
maintenance,
as a service

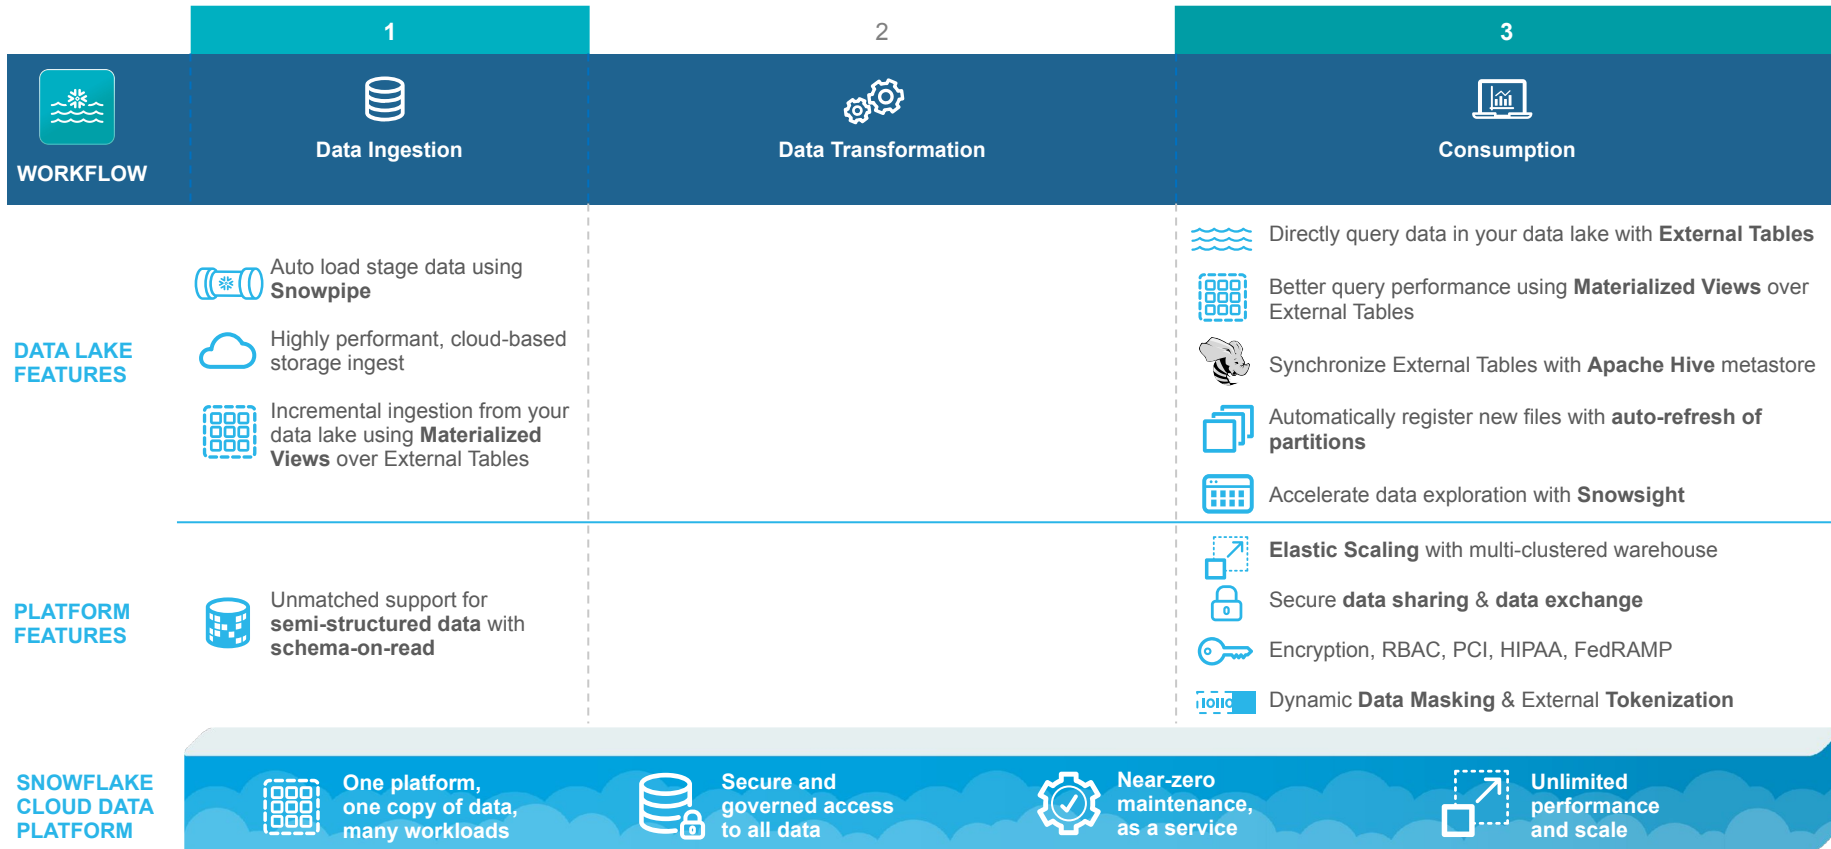


Unlimited
performance
and scale



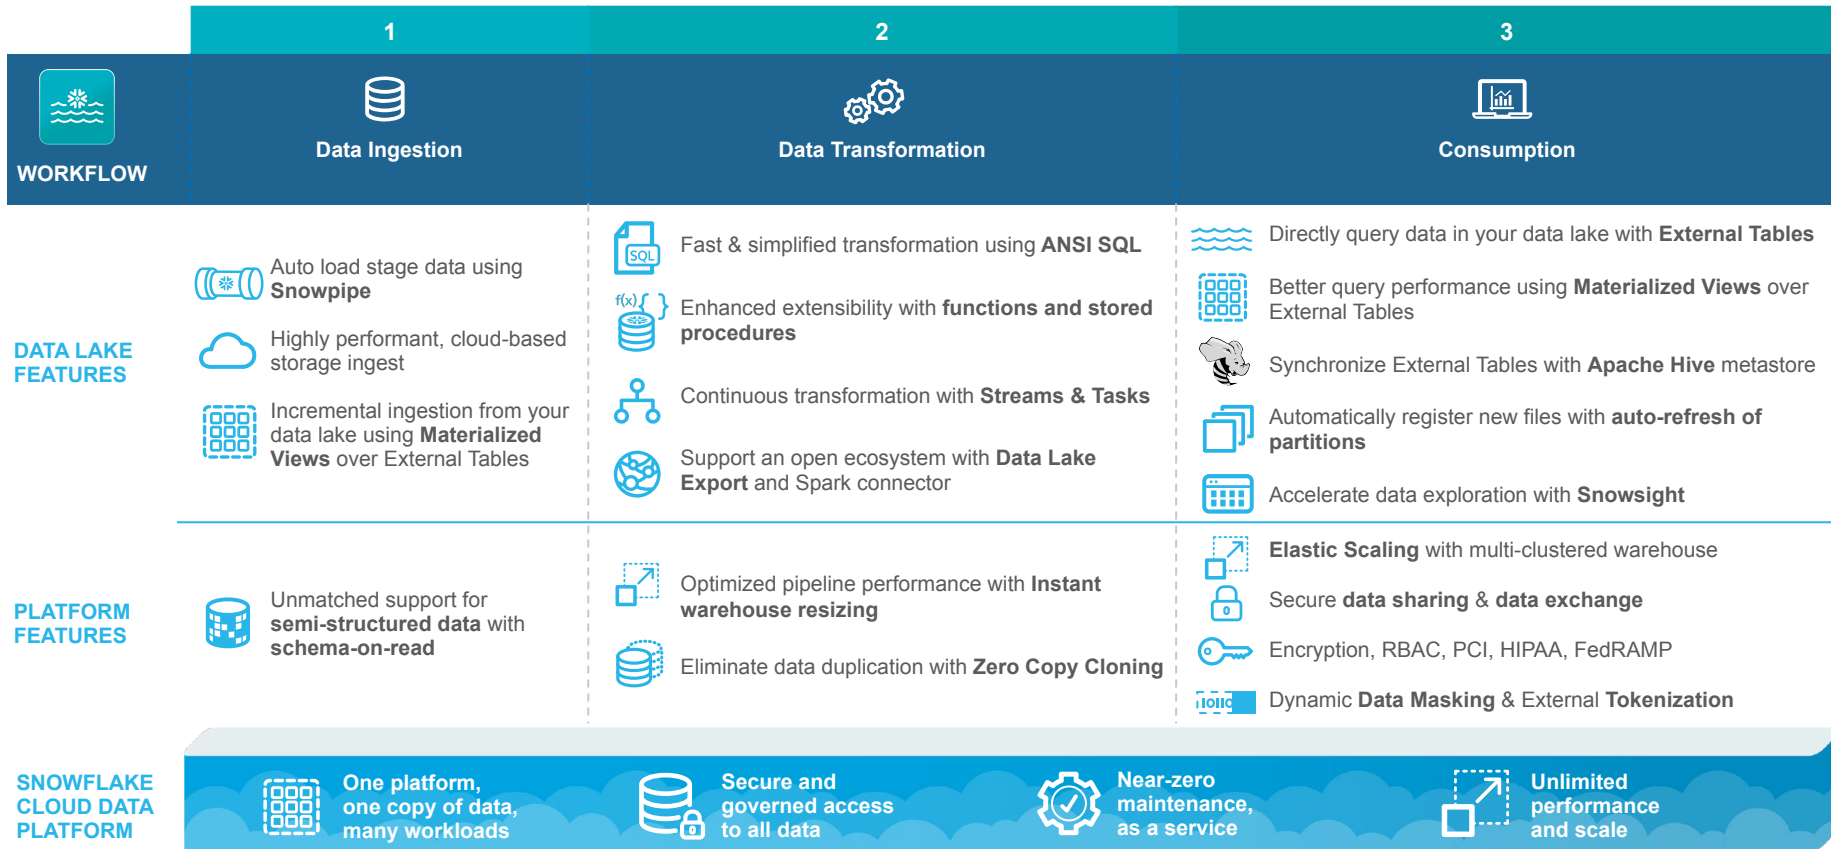
DATA LAKE WITH SNOWFLAKE

CLICK ON NAV BAR TO GO TO DEEP DIVE




DATA LAKE WITH SNOWFLAKE

CLICK ON NAV BAR TO GO TO DEEP DIVE



EXTERNAL TABLES



Use Case

Read-only tables over files on external cloud storage. Queries over archived or historical cloud data, ad-hoc queries over files.

Cost


Small maintenance overhead charged for manually refreshing the external table metadata (using `ALTER EXTERNAL TABLE ... REFRESH`).

Offset by Benefits

Query and access all data from a single governed data platform.



AUTO-REFRESH EXTERNAL TABLE PARTITIONS



GCS



S3



Azure

Use Case

Refresh external table metadata automatically using your clouds notification service

- New files in the path are added to the table metadata
- Changes to files in the path are updated in the table metadata.
- Files no longer in the path are removed from the table metadata.

Cost

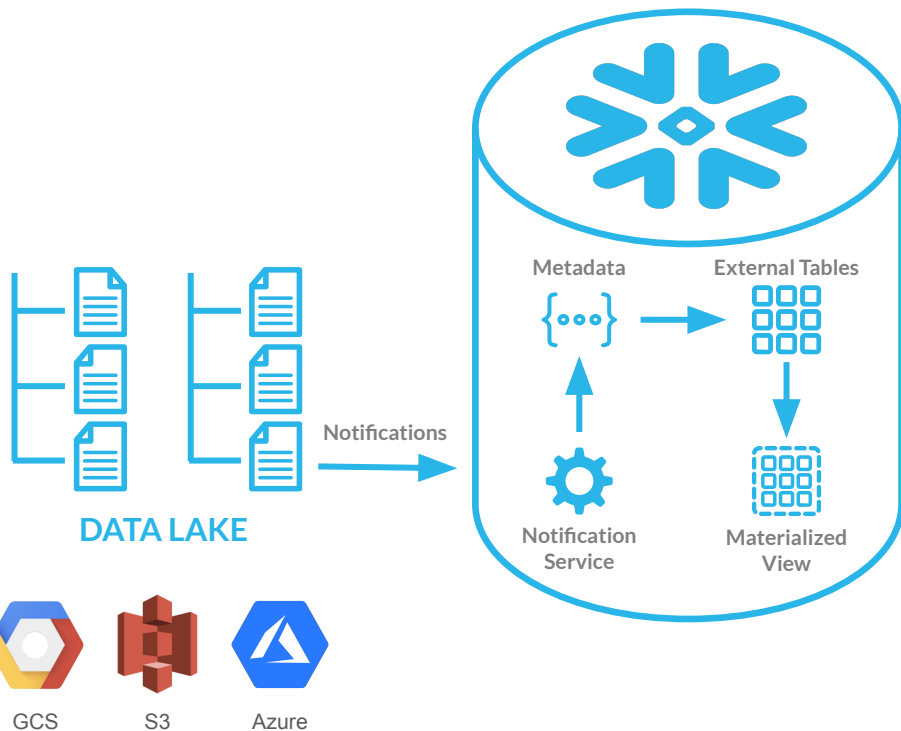
Small maintenance overhead charged for notification event management and refreshing table metadata.

Offset by Benefits

External data is always up to date and available to users and one location to access all data.



MATERIALIZED VIEWS



Use Case

Physically persist commonly aggregated data
MV on a Snowflake table guaranteed correct

Cost

Serverless refresh, additional storage

Offset by Benefits

Faster queries, fewer compute credits

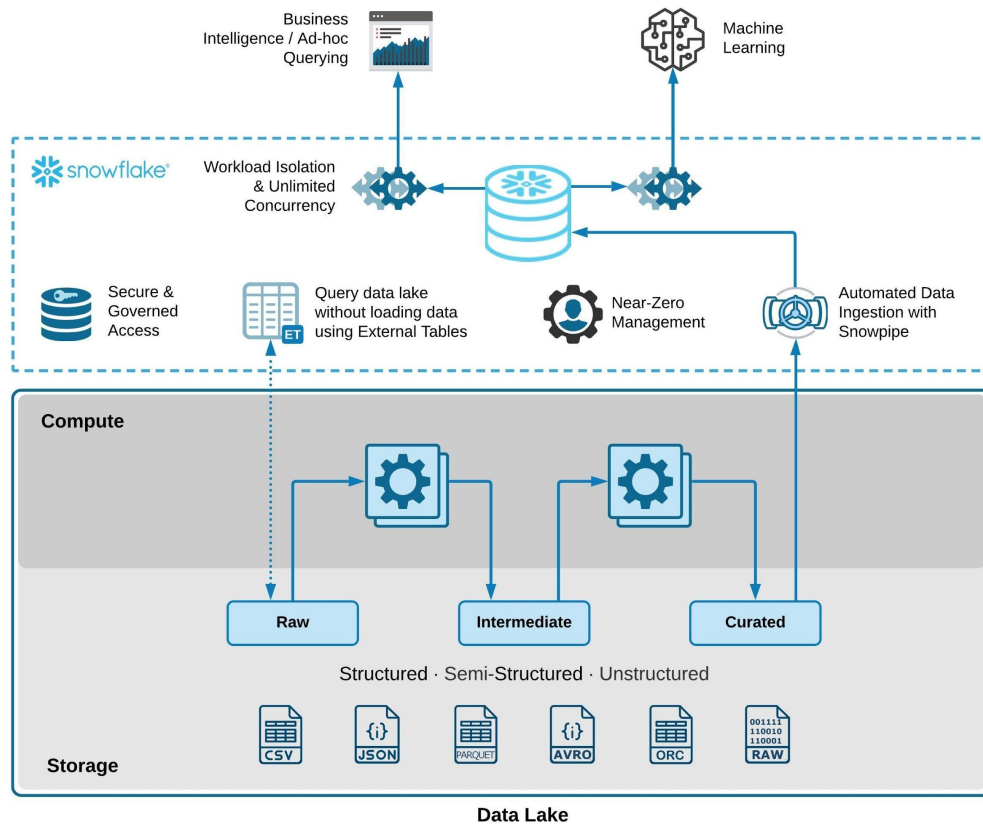
MV on External Tables

Fast, governed access to Data Lake *

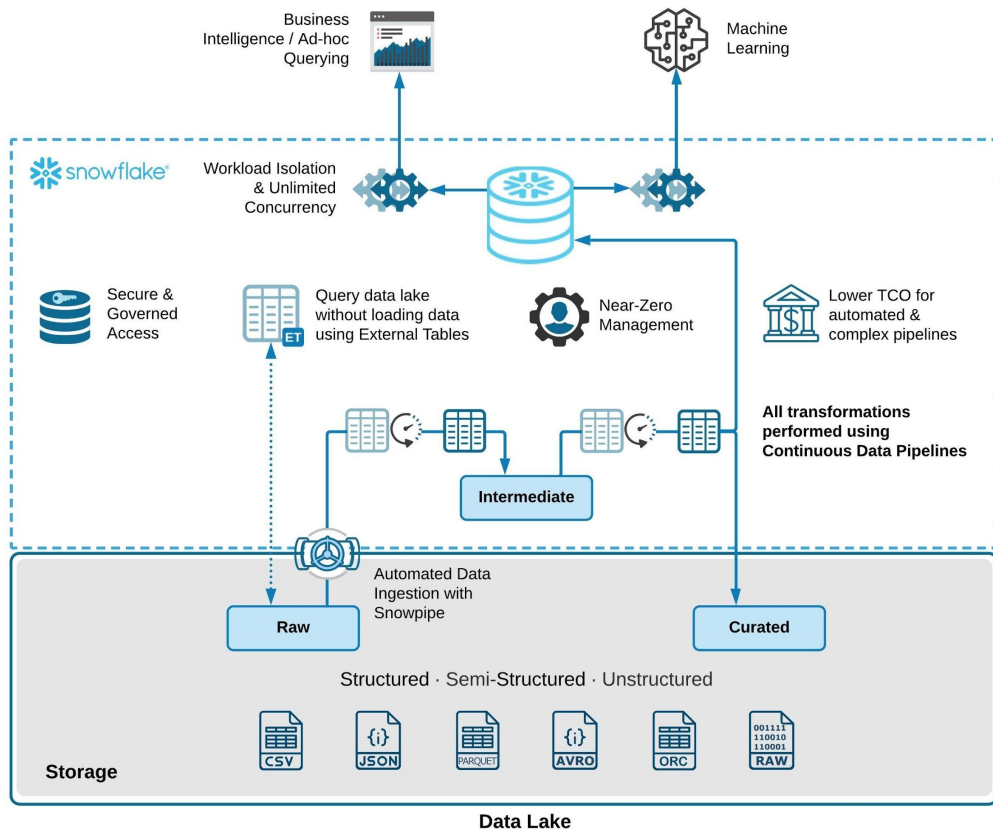
*Must automate / schedule a refresh



REFERENCE ARCHITECTURE #1: Snowflake as Fast Query Layer



REFERENCE ARCHITECTURE #2: Snowflake as Fast Query Layer + Transformation



Complement Your Data
Lake with Snowflake

Build Your Data Lake
with Snowflake



DATA LAKE WITH SNOWFLAKE

1



WORKFLOW



Data Ingestion

2



Data Transformation


3



Consumption



DATA LAKE WITH SNOWFLAKE



**SNOWFLAKE
CLOUD DATA
PLATFORM**



One platform,
one copy of data,
many workloads



Secure and
governed access
to all data



Near-zero
maintenance,
as a service



Unlimited
performance
and scale



DATA LAKE WITH SNOWFLAKE

CLICK ON NAV BAR TO GO TO DEEP DIVE

1

2

3



WORKFLOW



Data Ingestion



Data Transformation



Consumption

DATA LAKE
FEATURES



Auto load stage data using **Snowpipe**



Highly performant, cloud-based
storage ingest



Ingest streaming data using **Kafka
connector**



Storage compression

PLATFORM
FEATURES



Unmatched support for
semi-structured data with
schema-on-read

SNOWFLAKE
CLOUD DATA
PLATFORM



One platform,
one copy of data,
many workloads



Secure and
governed access
to all data



Near-zero
maintenance,
as a service

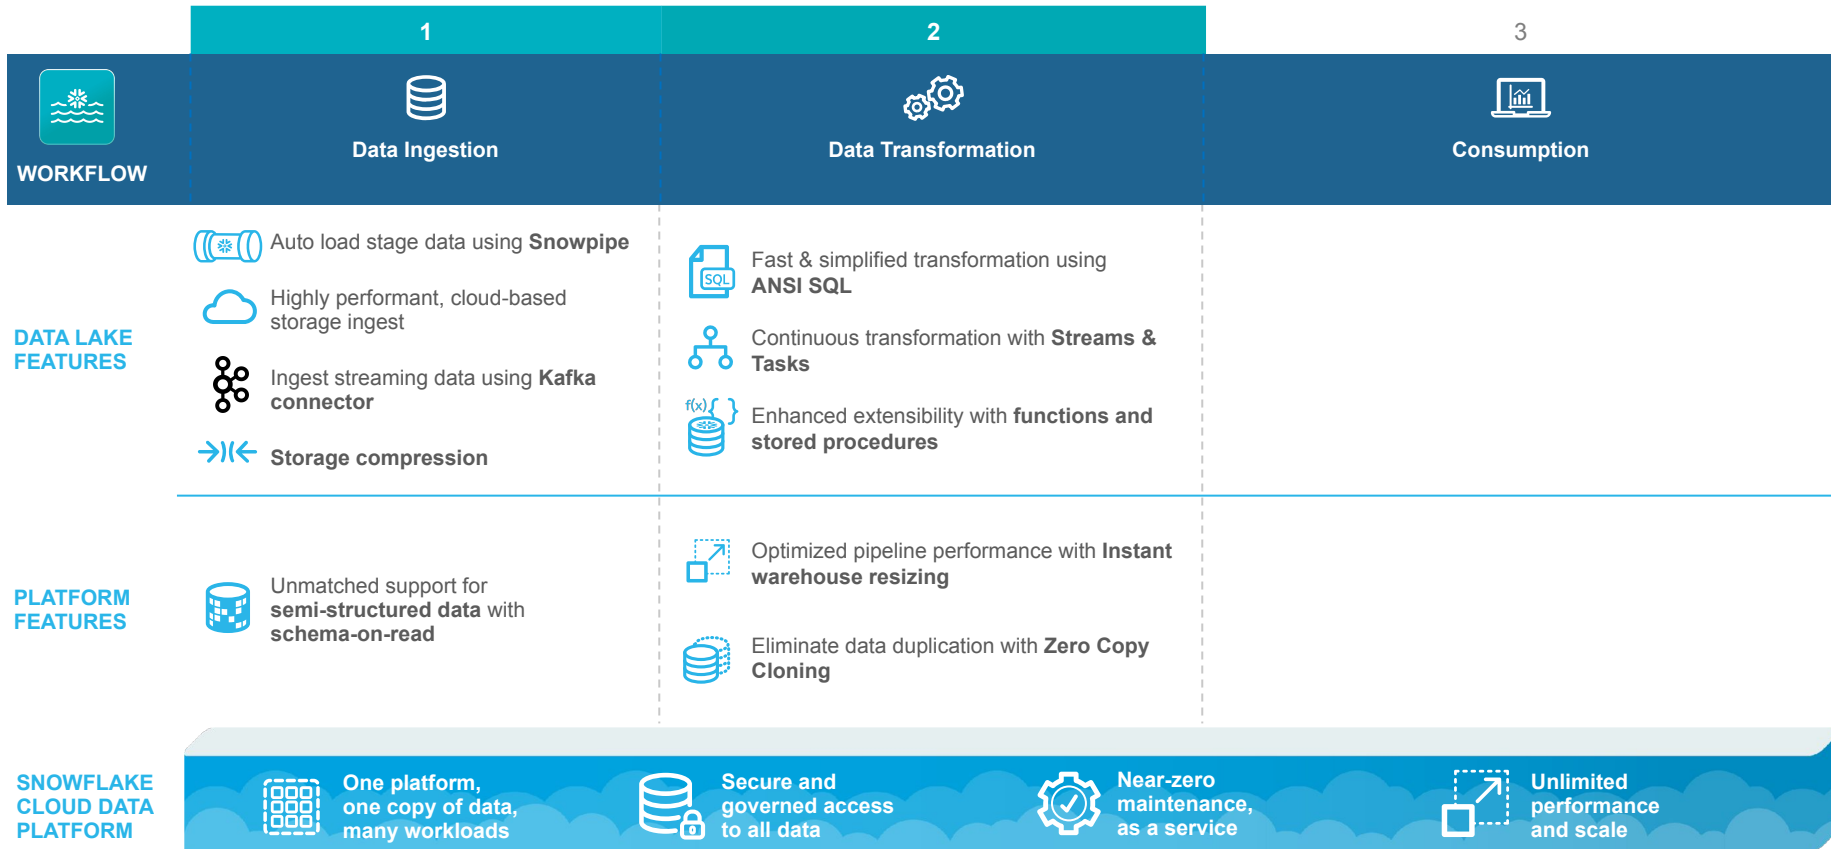


Unlimited
performance
and scale



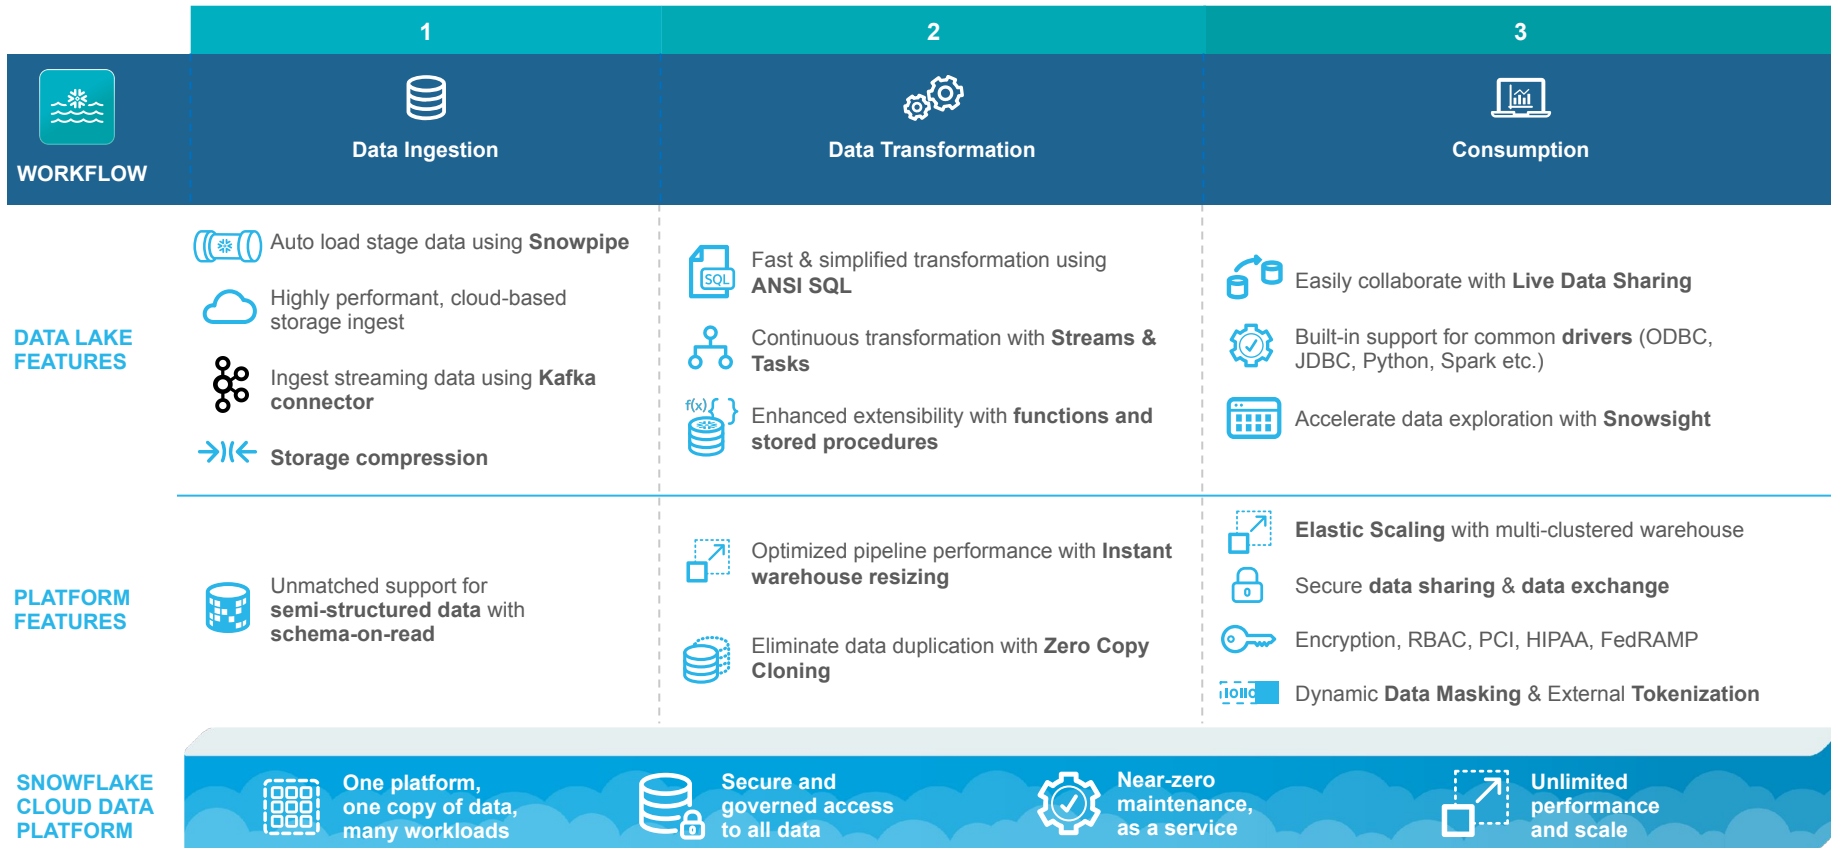
DATA LAKE WITH SNOWFLAKE

CLICK ON NAV BAR TO GO TO DEEP DIVE

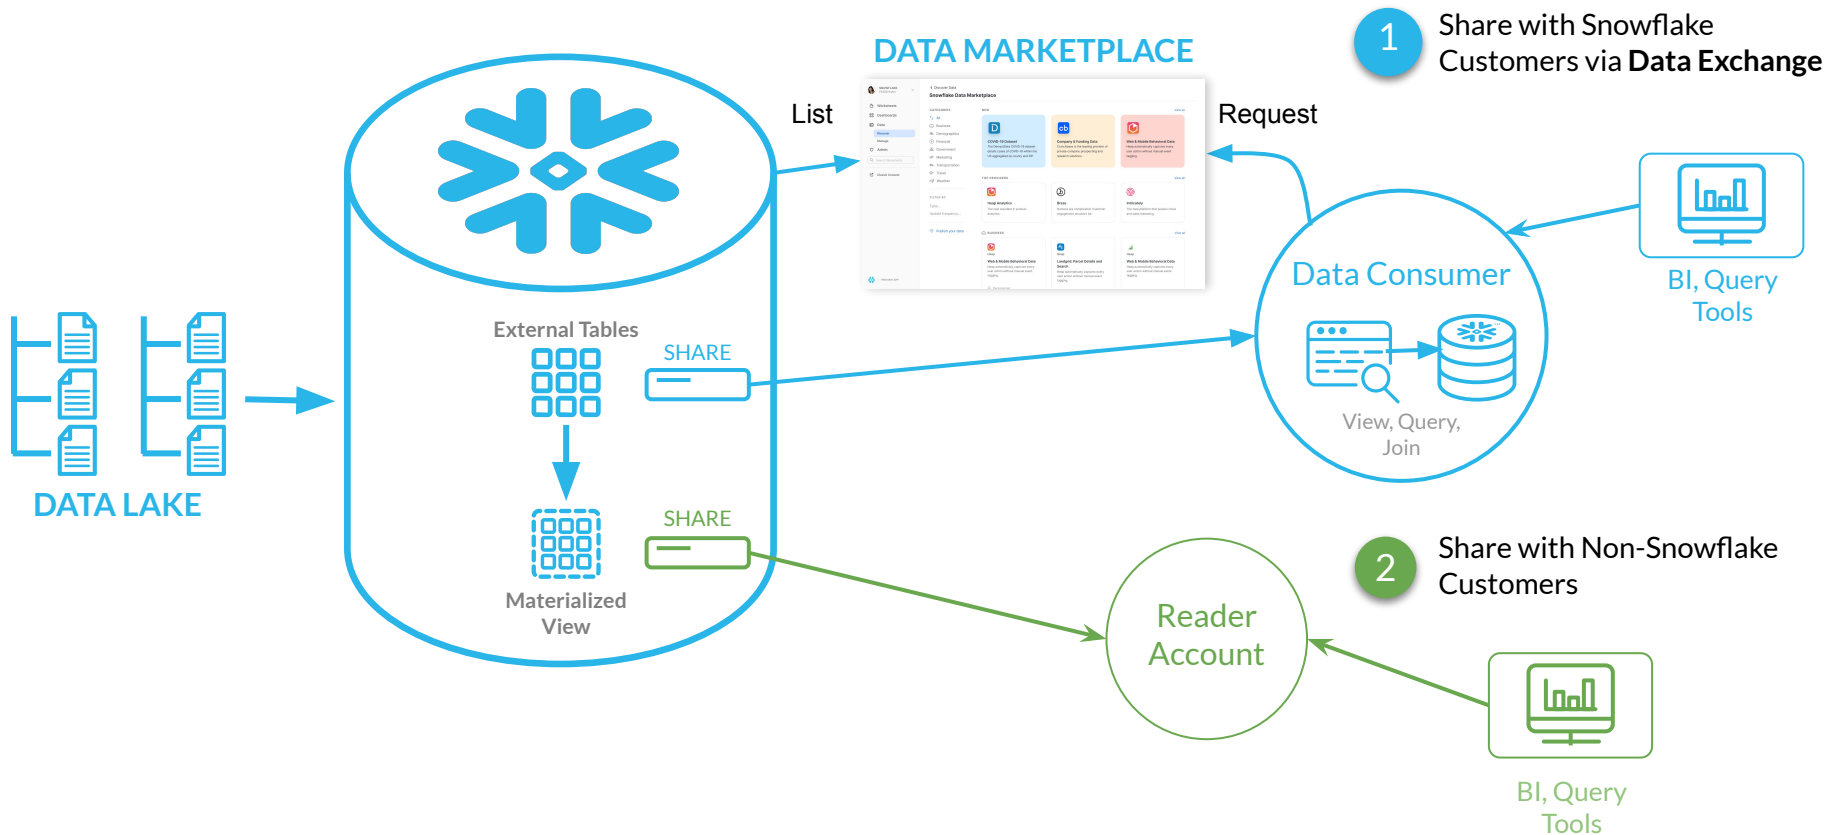


DATA LAKE WITH SNOWFLAKE

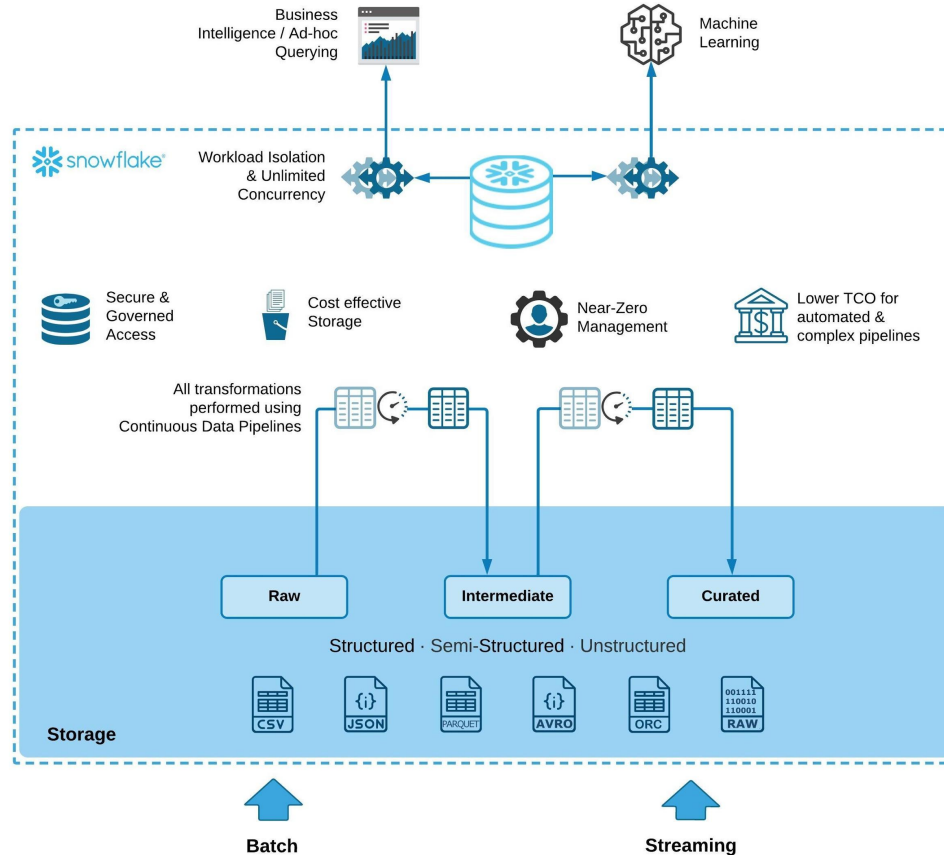
CLICK ON NAV BAR TO GO TO DEEP DIVE



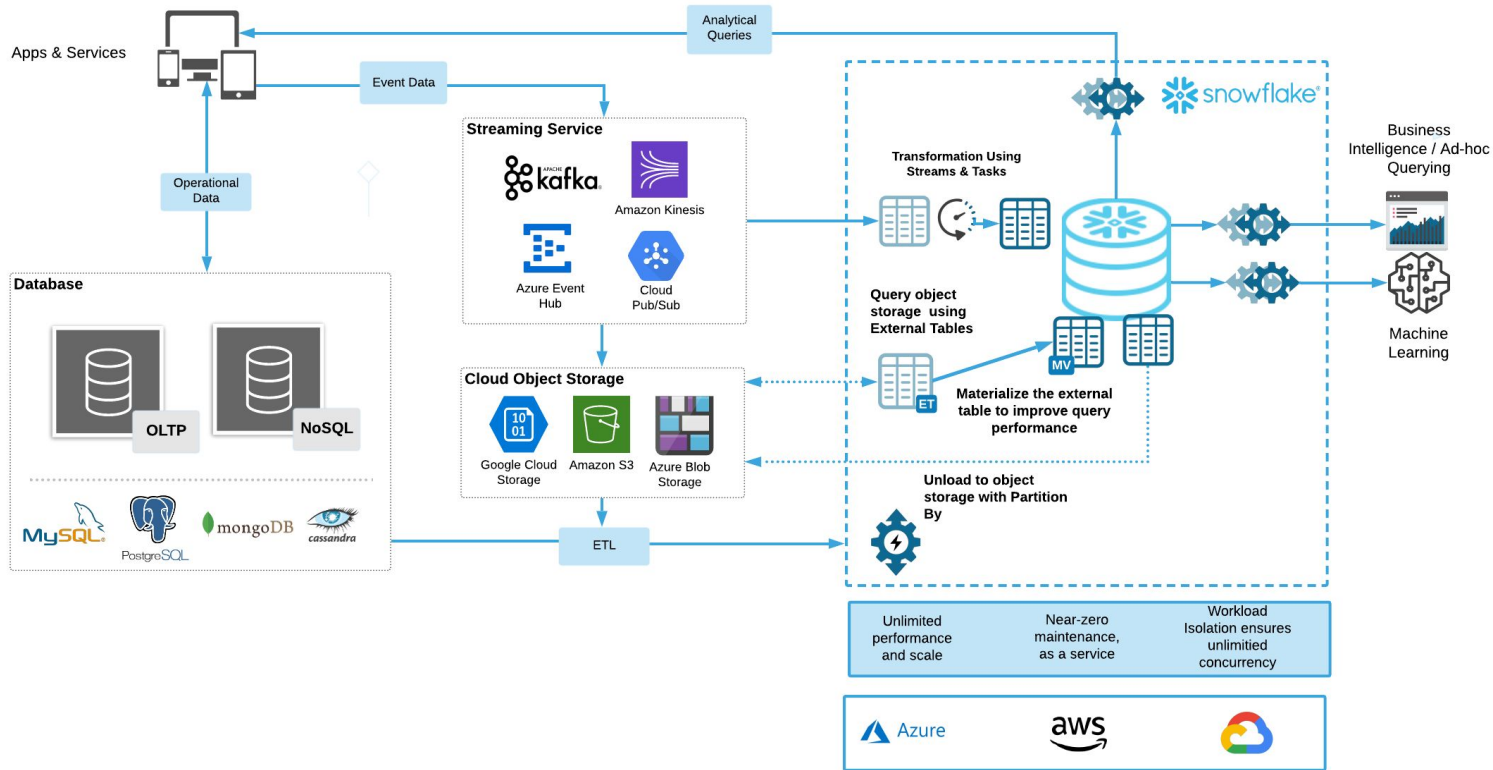
SHARE YOUR LAKE!



REFERENCE ARCHITECTURE #1: Snowflake Platform



REFERENCE ARCHITECTURE



SNOWFLAKE: SECURITY AT MULTIPLE LEVELS



Google Cloud

aws

Azure

You have security features built in to protect the data you load in and use in Snowflake

Snowflake uses the industry standard “shared responsibility” model

Snowflake uses many sophisticated mechanisms to keep the platform safe and stable

Orchestrated security built into the fabric of the platform, automated controls in place for all functions, constant monitoring and analysis to detect and mitigate threats quickly



PROTECTING YOUR DATA IN SNOWFLAKE

End-to-End Encryption

Always-encrypted client communications, plus integration with cloud provider private networking



Fully Encrypted Storage

Data at rest is always encrypted while handled by the Snowflake drivers and systems



Strong Authentication

Built in multi-factor, integration with your federated SSO, easy user management



Full Auditing

Track every login, every transaction, every data transfer, and export to your security tools



Role-Based Access Control

All objects, actions, and even compute usage can be controlled with roles



Recovery

We give you options to ensure your data can be recovered in case of an accident or worse



[Snowflake Security Product Documentation](#)



SNOWFLAKE WILL BE 100% TRANSPARENT

Third Party Attestations and Certifications

Snowflake Security and Trust Center: <https://www.snowflake.com/product/security-and-trust-center/>

Snowflake Security Policy: <https://www.snowflake.com/legal/> (first link)



SOC 2 Type II

12 Month Coverage Period

SOC 1 Type II

6 Month Coverage Period



PCI-DSS



ISO/IEC 27001



FedRAMP

(Available from OMB MAX)



HIPAA

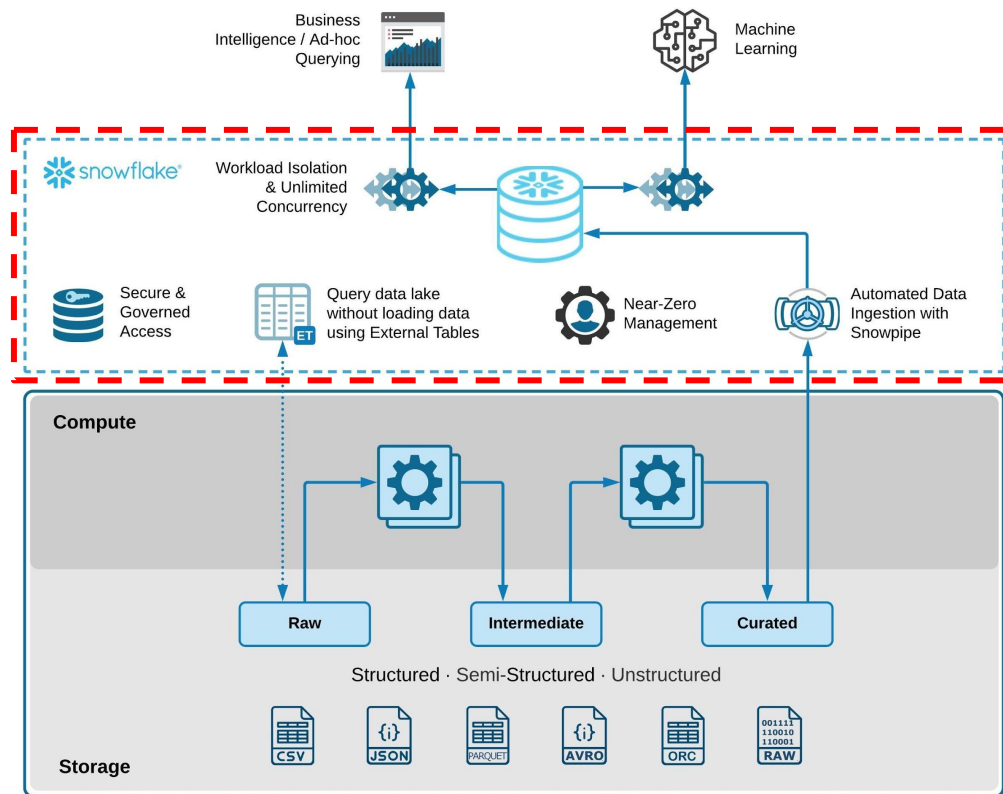
HITRUST Certified

Self-Assessments

CAIQ, SIG Lite, Pen Test Results



PCI, HIPAA, ISO 27001 etc.



PCI, HIPAA, ISO 27001 etc.

